

Where your code goes: **nowhere it didn't already.**

The under-claiming summary — what L1 produces, and what your org always owns.

Data flow // trust boundary

The developer workstation is the trust boundary. 3C adds no path off the workstation for source code, prompt content, or credentials, under any maturity level.

At L2+, the only thing 3C optionally emits is anonymised metadata — counts, pass/fail, durations — to a self-hostable backend. It does not change how your AI assistant already operates with its own model provider; that remains governed by your existing agreement with that vendor.

Security posture // by design

- **Local storage** — config + event state on the workstation (SQLite + JSONL); no code or prompt content off-box.
- **Audit trail** — append-only local log at L1; tamper-evident hash-chain + escrow signing is an L3 addition over the same path.
- **Offline resilient** — the workflow never blocks on a backend; no always-on egress.
- **White-label** — your org is the AIMS attribution holder, not the upstream framework.

What 3C does NOT do // honest

3C does **not** certify any compliance standard. It does **not** provide a complete regulatory evidence pack at any maturity level.

L1 produces evidence **substrate**, not evidence packs. Your org always owns the AI risk register, the DPIA/FRIA process, reviewer training, and the record-custodian designation.

Evaluate it as OSS // not a vendor

- **Licence:** MIT — read it in the repo [LICENSE](#).
- **SBOM:** `npm sbom --sbom-format cyclonedx` after install captures 3C + transitive deps; git-installed, so recorded with `git+https` provenance.
- **Source of truth:** every claim is verifiable in [kgn-git/praise](#) — the full requirements baseline is maintained there.