

Your team's AI, finding the same page.

Open, local-first, audit-ready — the one-page case for a calm rollout.

The moment this is for // why now

AI assistants made individual developers much faster without making teams more coherent. Each session is **locally optimal** — passing local tests, matching the file it can see — but the aggregate drifts: divergent conventions, drifting architecture, no evidence trail.

There is also gentle pressure to be able to *show your work* — which changes were AI-assisted, against which conventions, reviewed by whom. Teams that adopted AI with no alignment have no substrate to produce that from.

A path, walked at your pace // L1 → L4

A set of places a team can choose to be — each answering a real question at a clearly-stated commitment level. Stopping anywhere is a valid decision:

- **L1 Standardise** — shipped. Reversible in one PR. Eases architectural drift.
- **L2 Automate** — gates run automatically. Eases review-burden inflation.
- **L3 Observe** — evidence packs; mapped before the 2027-08 deadline.
- **L4 Optimise** — complexity-adapted governance. A horizon, not a date.

What L1 asks of you // honest

- **Adoption:** about a week from `npx 3c init`; starting point + delta measured in-week.
- **Infrastructure:** none — runs on developer workstations.
- **Lock-in:** none — Agent Skills open standard, MIT.
- **Maintenance:** a steward role ~2 hr/week per active 6–15-dev team.

What L1 gives back // adoption week

Visible by the end of the one-week adoption sprint, then compounding:

- Two engineers on one repo produce structurally consistent code without a reviewer forcing it.
- Issue → implement → review → test runs the same way on every laptop; onboarding is one `CLAUDE.md`.
- Lint, types, tests, secret-scan run before code review, every commit.

3C produces evidence **substrate**, not certification: it does not issue any compliance certification or conformity statement — your org owns the risk register.